

Machine Learning V

Non-Linear Models – Part II




Outline


- **Part I : The Need for non-linear models**
- **Part I : Extending the simple linear classifier**
 - Adding non-linearity
 - Simple Neural Networks
- **Part II : Support Vector Machines**
 - Linear SVMs
 - The “Kernel-Trick”

- Invented in the mid 90s by Vapnik
- Classification and Regression
- State of the Art ML Algorithm of the pre Deep Learning era

- Invented in the mid 90s by Vapnik
- Classification and Regression
- State of the Art ML Algorithm of the pre Deep Learning era
- **Basic model:**
 - Support only two classes $\{-1,1\}$
 - Linear classification




- Invented in the mid 90s by Vapnik
- Classification and Regression
- State of the Art ML Algorithm of the pre Deep Learning era
- **Basic model:**
 - Support only two classes $\{-1,1\}$
 - Linear classification



Parameterization: $wx - b = 0$

What is the difference compared to previous formulations?




Standard linear model:

- loss only on accuracy
- many solutions


$$\arg \min_w \sum_{i=0}^N L(y_i, w^T x_i + b)$$

What is the difference compared to previous formulations?



Standard linear model:

- loss only on accuracy
- many solutions



New optimization problem

- “Max Margin”: $\frac{2}{\|w\|}$

- only one solution, convex optimization problem

Basic formulation


New optimization problem

- maximize “Margin”:
- equals minimizing the uncertainty

$$\arg \min_w \frac{1}{2} \|w\|$$

subject to


$$y_i(\langle \mathbf{w}, \mathbf{x}_i \rangle + b) \geq 1$$



Soft Margin: allow some outliers

Still

- maximize “Margin”:
- equals minimizing the uncertainty
- enable linear separability



Soft Margin: allow some outliers


Still

- maximize “Margin”:
- equals minimizing the uncertainty
- enable linear separability

$$\arg \min_w \frac{1}{2} \|w\|^2 + C \sum_i \xi_i$$

subject to

$$y_i(\langle w, x_i \rangle + b) \geq 1 - \xi_i$$



Error

$$\xi_i = \max(0, 1 - y_i(w \cdot x_i - b))$$


Penalty factor

Dual Formulation of the optimization Problem

Via *Lagrange dual function*, leads to quadratic optimization problem (**convex!**)

Re-write w as linear combination of samples:

$$\mathbf{w} = \sum_{i=1}^m \alpha_i y_i \mathbf{x}_i$$



There is exactly one solution
and we will find it!

Dual Formulation of the optimization Problem


Via **Lagrange dual function**, leads to quadratic optimization problem (convex)

Re-write w as linear combination of samples:

$$w = \sum_{i=1}^m \alpha_i y_i x_i$$

$$\arg \max_{\alpha} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y_i y_j \langle x_i, x_j \rangle$$

Subject to $0 \leq \alpha_i \leq C$ and $\sum_{i=1}^m \alpha_i y_i = 0$



Normal as linear combination of samples


Dual Formulation of the optimization Problem

Via **Lagrange dual function**, leads to quadratic optimization problem (convex)

Re-write w as linear combination of samples:

$$w = \sum_{i=1}^m \alpha_i y_i x_i$$

$$\arg \max_{\alpha} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y_i y_j \langle x_i, x_j \rangle$$



Subject to $0 \leq \alpha_i \leq C$ and $\sum_{i=1}^m \alpha_i y_i = 0$

Dot-product over all samples

Dual Formulation of the optimization Problem


Via **Lagrange dual function**, leads to quadratic optimization problem (convex)

Re-write w as linear combination of samples:

$$w = \sum_{i=1}^m \alpha_i y_i x_i$$

$$\arg \max_{\alpha} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y_i y_j \langle x_i, x_j \rangle$$

Subject to $0 \leq \alpha_i \leq C$ and $\sum_{i=1}^m \alpha_i y_i = 0$



Soft Margin


Dual Formulation of the optimization Problem

Via **Lagrange dual function**, leads to quadratic optimization problem (convex)

Re-write w as linear combination of samples:

$$\mathbf{w} = \sum_{i=1}^m \alpha_i y_i \mathbf{x}_i$$

$$\arg \max_{\alpha} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y_i y_j \langle \mathbf{x}_i, \mathbf{x}_j \rangle$$



Subject to $0 \leq \alpha_i \leq C$ and

$$\sum_{i=1}^m \alpha_i y_i = 0$$


Weight balance between classes

Dual Formulation of the optimization Problem

Via *Lagrange dual function*, leads to quadratic optimization problem (convex)

Classification:

$$f(\mathbf{x}) = \text{sgn}(\langle \mathbf{w}, \mathbf{x} \rangle + b) = \text{sgn}\left(\sum_{i=1}^m \alpha_i y_i \langle \mathbf{x}_i, \mathbf{x} \rangle + b\right)$$




Dual Formulation of the optimization Problem

Via *Lagrange dual function*, leads to quadratic optimization problem (convex)

Classification:

$$f(\mathbf{x}) = \text{sgn}(\langle \mathbf{w}, \mathbf{x} \rangle + b) = \text{sgn}\left(\sum_{i=1}^m \alpha_i y_i \langle \mathbf{x}_i, \mathbf{x} \rangle + b\right)$$



Only a few vectors have $\alpha_i > 0$

→ “**Support Vectors**” form the actual model

Non Linear SVMs:

→ follow same strategy as before and add simple non-linear function

At the formulation of the model normal: $\phi: \mathbb{R}^{d_1} \rightarrow \mathbb{R}^{d_2}, \mathbf{x} \mapsto \phi(\mathbf{x})$

$$f(\mathbf{x}) = \text{sgn}(\langle \mathbf{w}, \mathbf{x} \rangle + b) = \text{sgn}\left(\sum_{i=1}^m \alpha_i y_i \langle \mathbf{x}_i, \mathbf{x} \rangle + b\right)$$



$$f(\mathbf{x}) = \text{sgn}(\langle \mathbf{w}, \phi(\mathbf{x}) \rangle + b) = \text{sgn}\left(\sum_{i=1}^m \alpha_i y_i k(\mathbf{x}_i, \mathbf{x}) + b\right)$$

Kernel function

$$k(\mathbf{x}_i, \mathbf{x}_j) = \langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle$$

“Kernel Trick”: replace explicit non-linear function by **kernel**

Popular Kernels:

Polynomial (of degree d)

$$k(x, y) = \langle x, y \rangle$$

Gauss (or RBF)

$$k(x, y) = \exp\left(-\frac{\|x-y\|^2}{2\sigma^2}\right)$$


Linear

$$k(x, y) = \langle x, y \rangle$$


$$f(\mathbf{x}) = \text{sgn}(\langle \mathbf{w}, \phi(\mathbf{x}) \rangle + b) = \text{sgn}\left(\sum_{i=1}^m \alpha_i y_i k(\mathbf{x}_i, \mathbf{x}) + b\right)$$

Training (optimization problem) and inference do not change!

Multi class problems:




Multi class problems: 1-vs-Rest



N models, take best.

Multi class problems: 1-vs-1



$N(N-1)/2$ models – tree execution, take best.

Lab exercises coming up ...